
Vlocalhost.AI

Installation Guide

Meeting notes that never leave your machine. This guide covers Windows, macOS and Linux, what to do when something goes wrong, and how to set up the optional summary model.

Version	1.1.1
Updated	14 August 2026
Download	https://github.com/antigravitysoham-eng/vlocalhost-core/releases/latest
Support	https://antigravitysoham-eng.github.io/vlocalhost-ai/support/
Licence	AGPL-3.0 — free forever, source open

What you need

- **Nothing but the download.** The installer carries its own Python and every library it needs. You do not need to install Python, pip, or anything else first.
- **About 400 MB of free disk space**, plus ~150 MB the first time you record, when the speech model downloads.
- **A microphone.** Any built-in one is fine.
- **Optional: Ollama**, only if you want written summaries. Without it you still get a complete, timestamped transcript of every meeting.

No account, no sign-in, no internet connection after the download. Audio and transcripts never leave your computer.

Windows

Windows 10 or 11, 64-bit.

Install

1. Download **vlocalhost-1.1.1-windows-x64-setup.exe** from the downloads page.
2. Double-click it.
3. **Windows will show a blue “Windows protected your PC” screen.** This is expected — see the note below. Click **More info**, then **Run anyway**.
4. Choose where to install it. The default is inside your user folder and needs no administrator rights. You can pick any drive — `D:\Apps\Vlocalhost` is fine.
5. Leave **Create a desktop shortcut** ticked if you want one.
6. Click **Install**, then **Finish**. The app opens and asks you two setup questions.

Why the warning appears. Windows shows that screen for any program without a paid code-signing certificate. It is not a virus warning and says nothing about the file's contents. Vlocalhost is open source — you can read every line of what the installer puts on your machine. A signed build is planned.

The installer needs no administrator password, so it works on a locked-down work laptop.

Uninstall

Settings → **Apps** → **Installed apps** → **Vlocalhost.AI** → **Uninstall**, or use the shortcut in the Start menu folder. Your meeting notes and settings are deliberately kept. To remove those too, delete this folder afterwards:

```
%LOCALAPPDATA%\Vlocalhost
```

Or: the portable ZIP

Every release also ships **vlocalhost-1.1.1-windows-x64.zip**: the same build with no installer. Take this one if a virus scanner refuses the .exe, if Chrome reports “Virus scan failed”, or if you simply want nothing written outside a single folder.

1. Unpack the .zip anywhere you can write. Keep the path short — a folder buried deeper than about 140 characters cannot unpack cleanly, because the files inside are already long.
2. Open the folder and run **Vlocalhost.cmd**.
3. That is the whole install. Nothing is written outside this folder.

No desktop icon? Correct — the ZIP deliberately touches nothing outside its own folder. The app can still make one. Run this once from the unpacked folder:

```
runtime\pythonw.exe app\vlocalhost.py --install-shortcut
```

It adds Vlocalhost.AI to your desktop and Start menu, pointing at the copy you unpacked, and finds the real desktop folder even when OneDrive has moved it.

Removing the portable version

There is no uninstaller and no Add/Remove Programs entry — that is what portable means. Removing it is one step, or three if you made shortcuts:

1. Delete the folder you unpacked. The program is gone.
2. If you created shortcuts, remove them with `runtime\pythonw.exe app\vlocalhost.py --remove-shortcut` before deleting the folder — or delete the two .lnk files by hand afterwards.
3. Notes and settings survive on purpose, in the folders listed under “Where your files live”. Delete those only if you want the meetings gone too.

macOS

Coming soon. Apple silicon and Intel builds are being packaged and tested. Everything below is how it will work; the download is not on the releases page yet. Today, both platforms run from source: see the README in the repository.

macOS 12 Monterey or newer.

Pick the right build

Click the Apple menu → **About This Mac** and look at the chip:

Your Mac says	Download
Apple M1 / M2 / M3 / M4	vlocalhost-1.1.1-macos-arm64.dmg
Intel	vlocalhost-1.1.1-macos-x64.dmg

Install

1. Download the .dmg and double-click it.
2. Drag **Vlocalhost** onto the **Applications** folder in the window that opens.
3. Eject the disk image, then open **Vlocalhost** from Applications or Launchpad.
4. **macOS will refuse the first launch** — see below.
5. The first time you record, macOS asks for **microphone permission**. Allow it, or the app records silence.

“Apple could not verify Vlocalhost”

macOS blocks apps that have not been notarised by Apple, which requires a paid developer account. Until that is in place, open it manually once:

1. Try to open the app. Dismiss the warning.
2. Open **System Settings** → **Privacy & Security**.
3. Scroll down. There is a line naming Vlocalhost with an **Open Anyway** button. Click it.
4. Confirm. macOS remembers, so you only do this once.

*On older macOS the shortcut is to right-click the app and choose **Open**. That no longer works on macOS 15 and later — use System Settings.*

Recording the other people on a call

macOS has no built-in way to capture what your speakers are playing, so out of the box Vlocalhost records only your microphone. To capture the far end, install **BlackHole** (free, open source), route your output through it, then choose it in **Settings** → **What to listen to**. Windows and Linux need no such step.

Uninstall

Drag **Vlocalhost** from Applications to the Bin. To remove notes and settings as well:

~/Library/Application Support/Vlocalhost

Linux

Coming soon. The Linux archive is being packaged and tested. Everything below is how it will work; the download is not on the releases page yet. Today, both platforms run from source: see the README in the repository.

64-bit, glibc 2.35 or newer (Ubuntu 22.04+, Debian 12+, Fedora 36+).

Install

Unpack the archive into your home folder and run the install script inside it. It installs for your user only, needs no root, and adds Vlocalhost.AI to the applications menu. The exact commands are published with the build, once the build has been tested on each distribution named above.

*The microphone needs **libportaudio2**. The script names the package for your distribution rather than guessing.*

Notes and settings are kept, in `~/ .local/share/vlocalhost`.

First run

The first time it opens, Vlocalhost asks two questions. Both can be changed later in **Settings** → **Models**, and the whole wizard can be re-run from **Settings** → **Run setup again**.

1. How accurate should transcription be?

Profile	Model	Memory	Best for
Light	tiny	~250 MB	Old or low-power machines, 2 cores
Balanced	base	~350 MB	The default. Keeps up with live speech
Accurate	small	~735 MB	4+ fast cores; noticeably better on Indian languages

You can also point it at your own model folder. The model downloads once, the first time you record — expect a wait of a minute or two on that first recording only.

2. Do you want written summaries?

Summaries are written by a second model running locally through **Ollama**. This is entirely optional:

- **Skip it** and you still get a complete timestamped transcript of every meeting, saved as a .txt file. Nothing else is lost.
- **Set it up** and you additionally get a Markdown file with a summary, key points, decisions and action items.

If Ollama is not installed, get it from **ollama.com**, then return to Settings and press **Run setup again**. The wizard detects it and offers to download the summary model (about 2 GB).

Where your files go

Notes and settings	
Windows	%LOCALAPPDATA%\Vlocalhost
macOS	~/Library/Application Support/Vlocalhost
Linux	~/.local/share/vlocalhost

These are deliberately outside the program folder, so updating, reinstalling or uninstalling never touches your meetings. Not sure? The app will tell you:

```
Vlocalhost --paths
```

Troubleshooting

If your problem is not here, the support page has a form and a way to send a diagnostic report:
<https://antigravitysoham-eng.github.io/vlocalhost-ai/support/>

Installing

Chrome: “Failed — Virus scan failed”

Not a detection. Chrome hands the finished file to Windows to scan, and that call failed.

Usually a third-party antivirus, a managed work laptop, or a Defender service that is not running. Press **Retry** in Downloads first. If it fails again, take the portable ZIP instead — same build, no installer — or verify the file against the SHA-256 published beside every download and keep it.

The installer stops part-way and rolls back

The folder you chose is too deep.

Windows limits a path to 260 characters, and the files inside the bundle are already long. An install folder beyond roughly 140 characters cannot fit. Nothing is left behind — install again somewhere shorter, such as the default.

Windows: “Windows protected your PC”

Expected. The installer is not code-signed yet.

Click **More info** → **Run anyway**. If the button is missing, your workplace may block unsigned installers — ask IT, or install from source.

Windows: antivirus removed the installer

Some scanners flag any unsigned installer that unpacks an interpreter.

Restore it from quarantine and add an exclusion, or verify the download first: each release publishes a SHA-256 checksum, and the .sha256 file next to the download should match what you compute locally.

macOS: “Vlocalhost is damaged and can't be opened”

Misleading message. It almost always means the quarantine flag is set, not that the file is corrupt.

Use System Settings → Privacy & Security → Open Anyway, as described in the macOS section. If that does not appear, clear the flag directly:

```
xattr -dr com.apple.quarantine /Applications/Vlocalhost.app
```

macOS: the app closes instantly

Usually the wrong build — an Intel bundle on Apple silicon or the reverse.

Check the chip under Apple menu → About This Mac and download the matching .dmg.

Linux: ./install.sh: Permission denied

The extracted script lost its executable bit.

```
bash install.sh
```

Linux: version `GLIBC_2.35' not found

Your distribution is older than the build.

Install from source instead — the repository README covers it, and it works on much older systems.

Starting

Nothing happens when I click the icon

The window failed to open and the error went nowhere, because the shortcut deliberately runs without a console.

Run it from a terminal to see the reason:

```
Windows: "%LOCALAPPDATA%\Programs\Vlocalhost\runtime\python.exe" \  
"%LOCALAPPDATA%\Programs\Vlocalhost\app\vlocalhost.py"  
  
Linux: ~/.local/opt/vlocalhost/runtime/bin/python3 \  
~/.local/opt/vlocalhost/app/vlocalhost.py
```

It opens in a terminal instead of a window

The graphical toolkit could not start, so it fell back.

On Linux this usually means no display is available over SSH. The terminal mode is fully functional — press Ctrl+C to stop and save.

Recording

The transcript is empty, or every line is silence

The microphone is muted, denied, or the wrong device.

Check the operating system's microphone permission first (macOS: System Settings → Privacy & Security → Microphone; Windows: Settings → Privacy → Microphone). Then confirm the device:

```
Vlocalhost --devices
```

I can hear the other people but they are not transcribed

A microphone only carries your half of a call.

Set **Settings** → **What to listen to** to **Both**. On macOS this additionally needs BlackHole — see the macOS section.

The transcript lags behind the meeting

The model is too large for the machine.

Switch to the **Light** profile in Settings → Performance, or press **Benchmark this machine** to measure what your hardware can actually keep up with.

Hindi (or another language) comes out as nonsense English

An English-only model was selected.

Models ending in .en are English only, and they hallucinate rather than fail. Choose a model without the suffix. The app blocks the combination where it can.

Summaries

The transcript saved but there is no notes file

Ollama was not reachable when the recording stopped.

The transcript is never lost to this — it is written first, on purpose. Check the server is running by opening this in a browser; you should see a short JSON response, not an error:

```
http://localhost:11434/api/tags
```

Then **Settings** → **Models** → **Re-check Ollama**.

Settings says the model is not installed, but I installed it

Ollama names models with a tag, so `llama3.2` and `llama3.2:latest` are different strings.

Press **List...** next to the summarisation box. It reads the models straight off the server, so you pick a name that certainly exists rather than typing one.

Getting help

Write a diagnostic report and attach it to your message. It lists your settings, versions and recent log lines — **no meeting audio, transcripts or notes**. Read it before sending; it is a plain text file.

```
Vlocalhost --diagnose
```

Support: <https://antigravitysoham-eng.github.io/vlocalhost-ai/support/>

Source and issues: <https://github.com/antigravitysoham-eng/vlocalhost-core>